

6.172
Performance
Engineering
of Software
Systems



LECTURE 1
**Introduction &
Matrix Multiplication**
Saman Amarasinghe
2014

**SPEED
LIMIT**



PER ORDER OF 6.172

Outline

Why performance engineering

6.172 administrivia

Case study: Matrix multiplication

WHY PERFORMANCE ENGINEERING?

SPEED
LIMIT

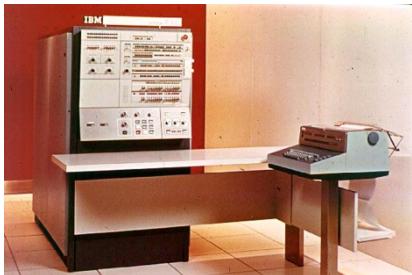


PER ORDER OF 6.172

In the Golden Era of Computing, Performance Engineering Ruled the World

Every programmer was a Performance Engineer

IBM System/360



Launched: 1964
Clock rate: 33 KHz
Data path: 32bits
Memory: 524 Kbytes
Cost: \$5,000 per month

Apple II



Launched: 1977
Clock rate: 1 MHz
Data path: 8 bits
Memory: 48 Kbytes
Cost: \$1,395

Any useful program will stretch the machine resources

Program has to be planned around the machine
Many will not ‘fit’ without intense performance hacks

Software Properties

What programmers want to do?

- New Functionality

... and...

- Scalability
- Compatibility
- Correctness
- Clarity
- Debuggability
- Maintainability
- Modularity
- Portability
- Reliability
- Robustness
- Testability
- Usability

... and more.

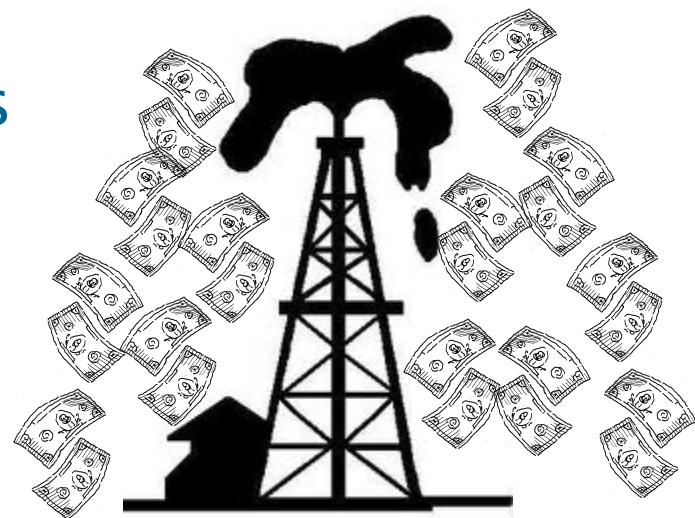
Performance is the **currency** of computing. You can often "buy" needed properties with performance.

In the Dominant Era of Computing, Performance became Free

The currency was free

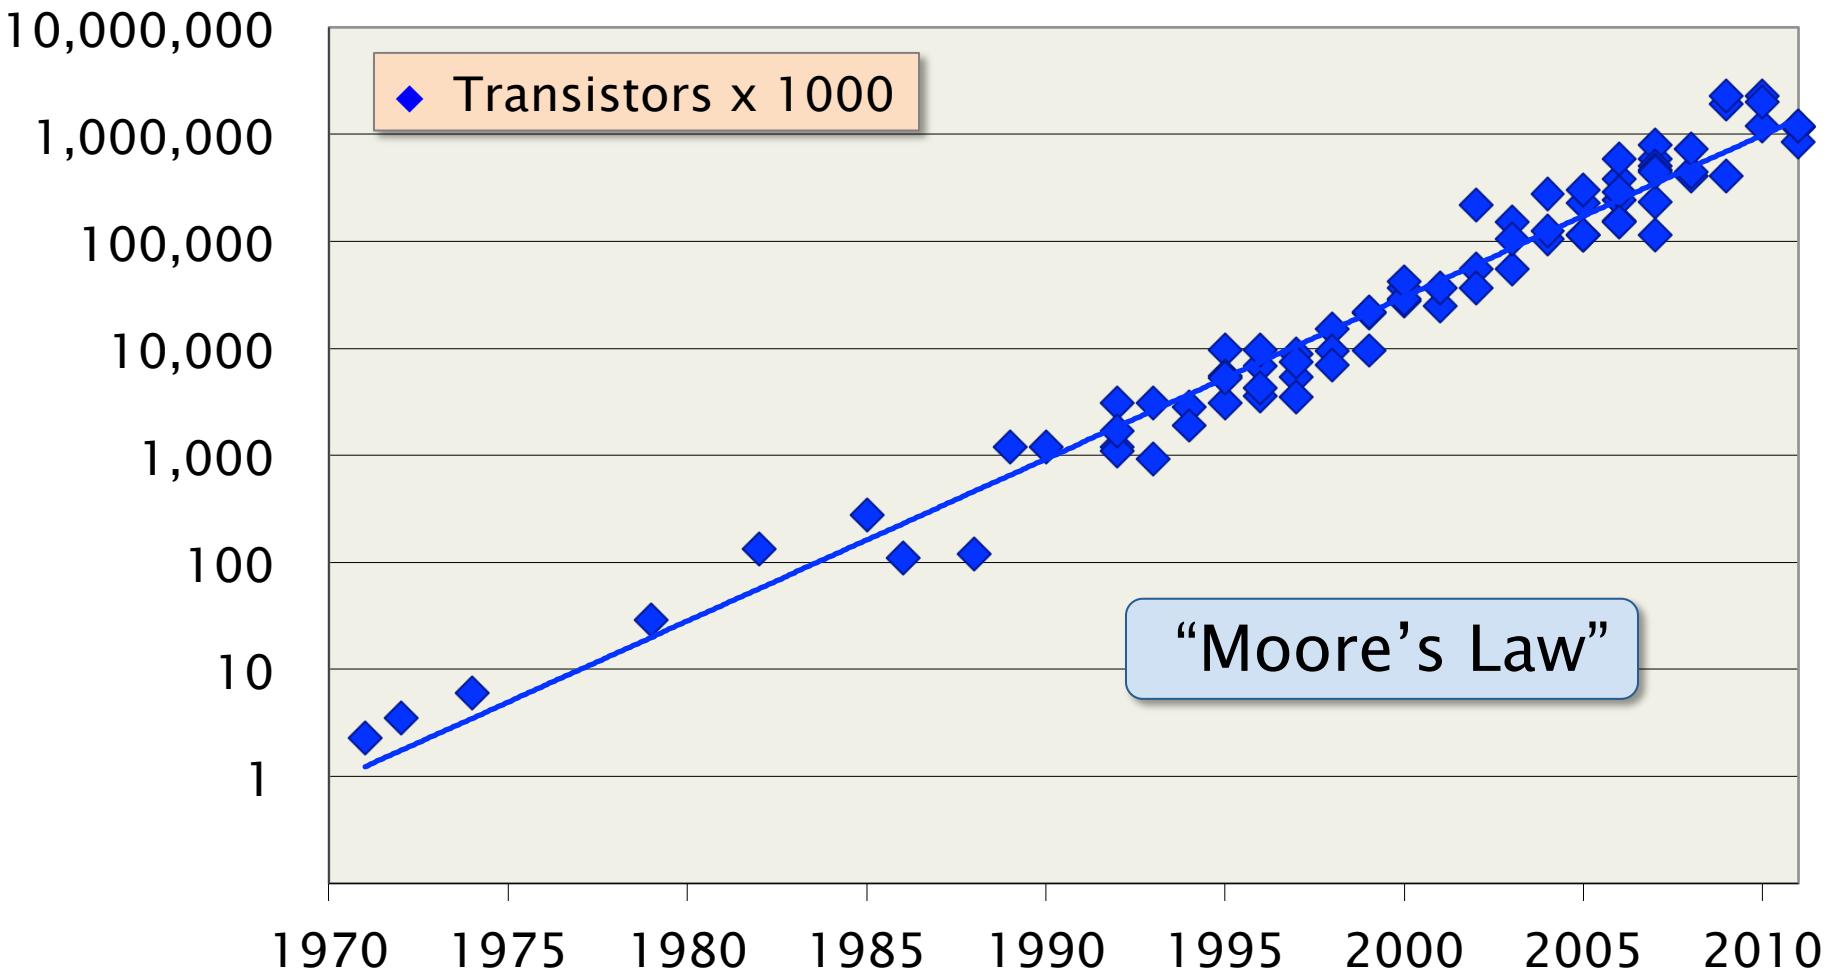
Only need to wait a few months

- Performance doubled every 2 years



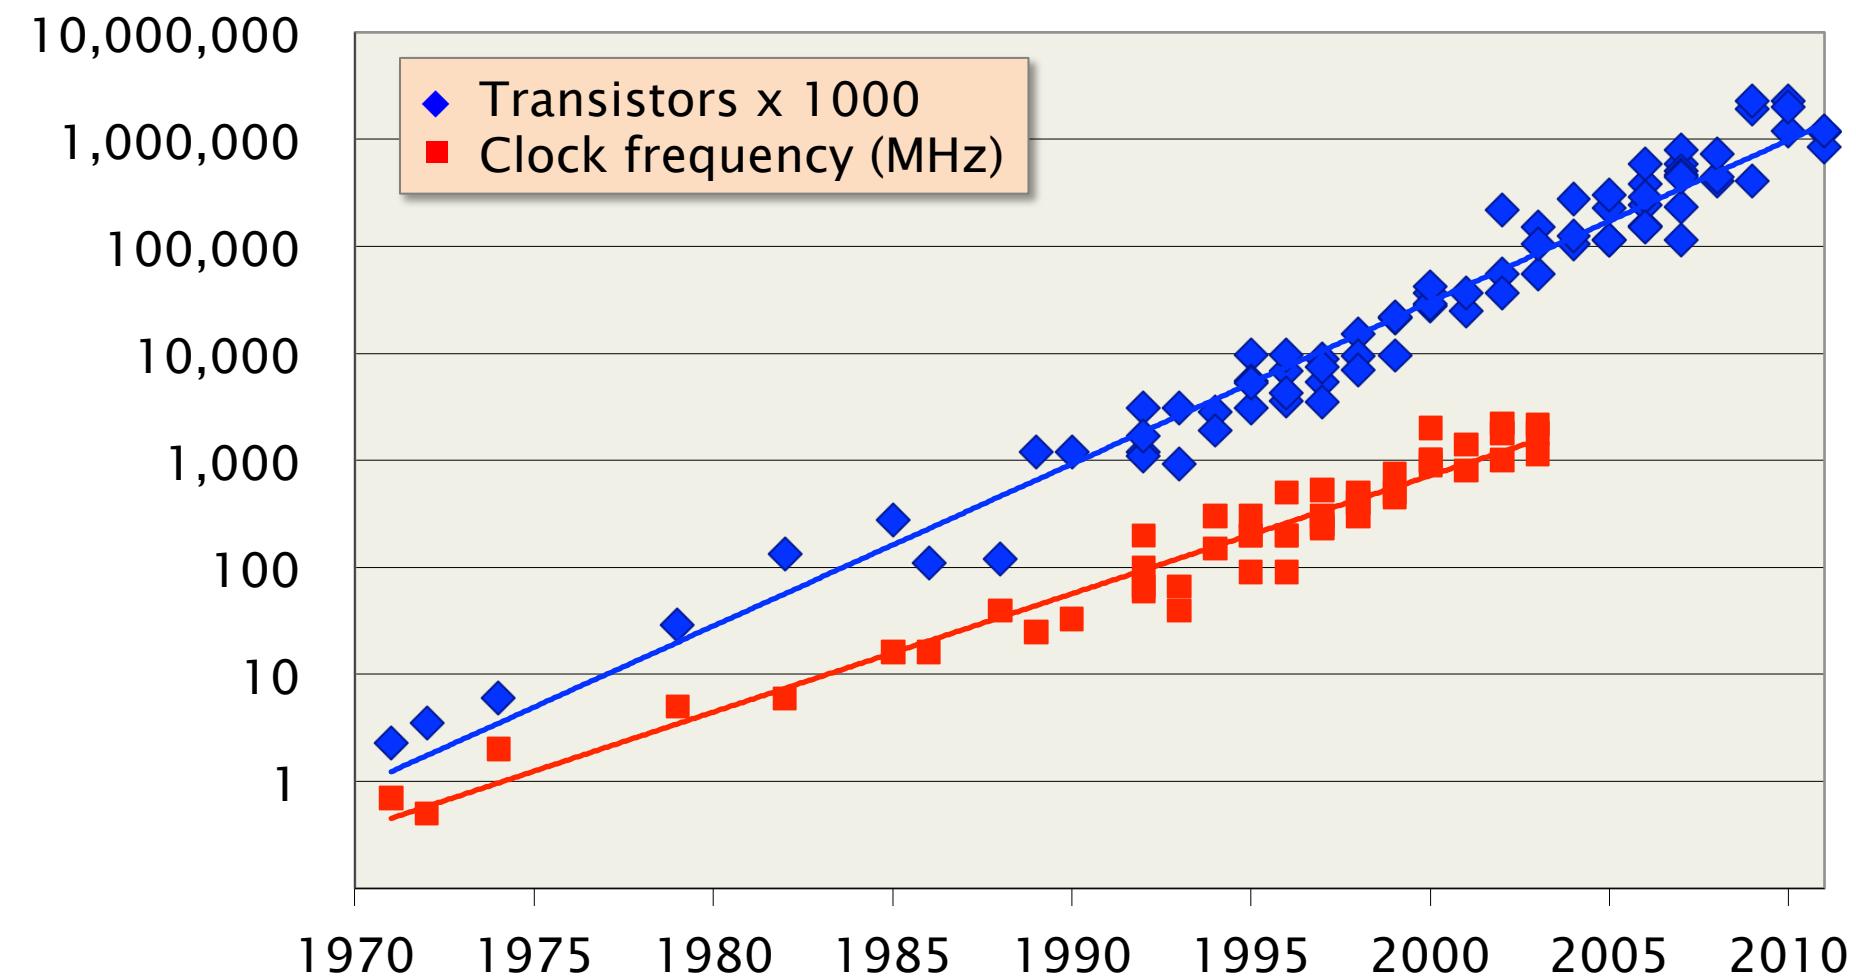
Performance is the **currency** of computing. You can often "buy" needed properties with performance.

Technology Scaling



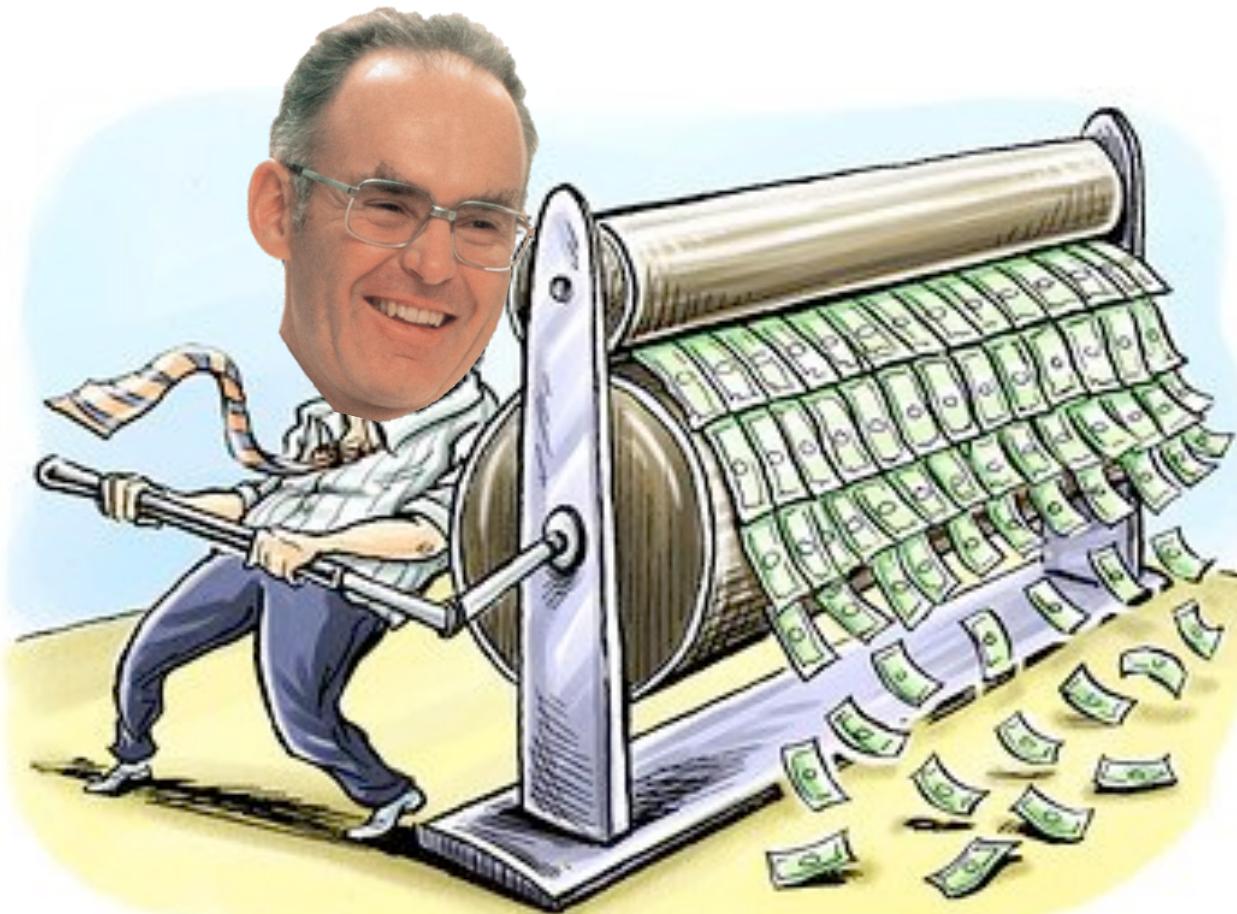
Intel processor chips

Technology Scaling



In the Dominant Era, Performance was Free

Moore's Law and the scaling of clock frequency
= printing press for the currency of performance



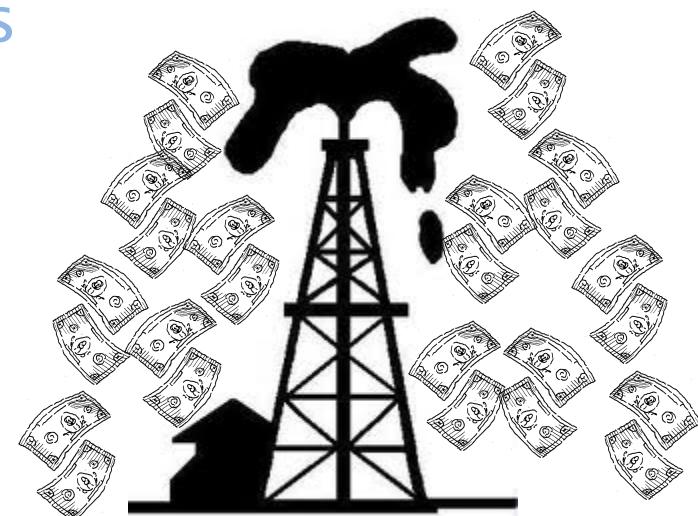
In the Dominant Era, Performance was Free

The currency was free

Only need to wait a few months

Performance doubled every 18 months

Performance engineering was
'optional' at best and
'irrelevant' in the eyes of most
programmers



Performance is the **currency** of computing. You can often "buy" needed properties with performance.

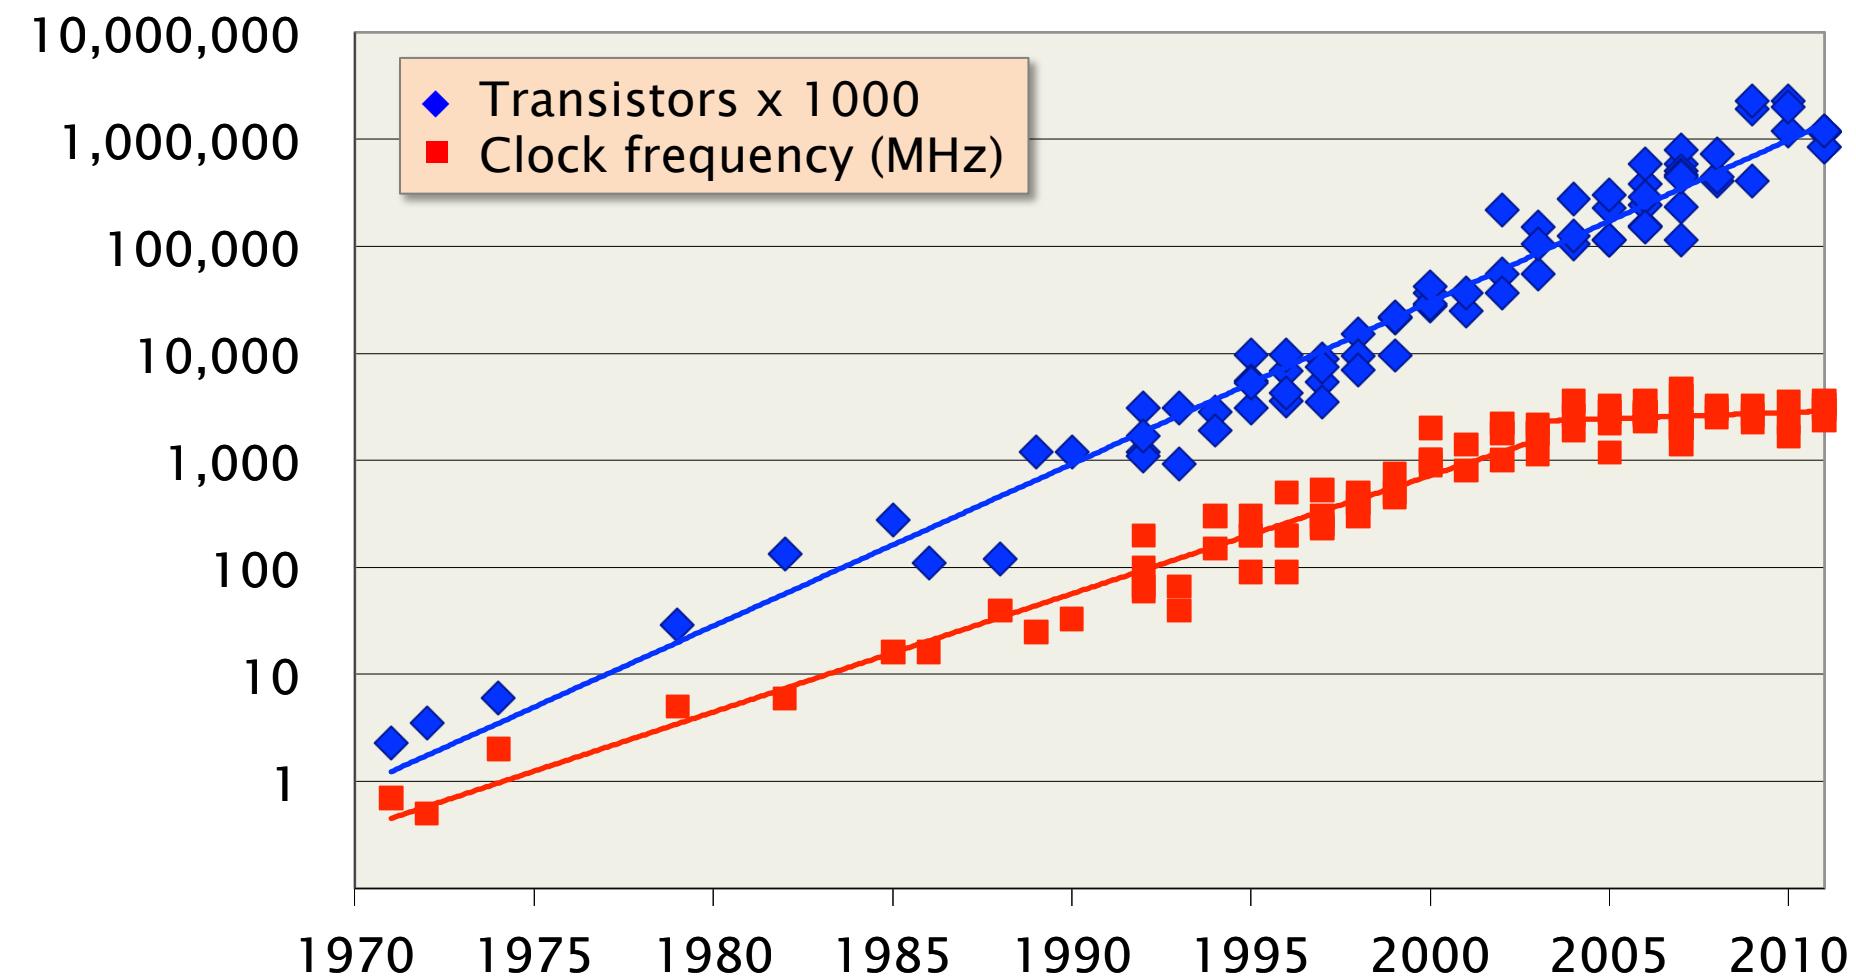
The Modern Era of Computing, Programmers need to Work for Performance

Moore's law is not giving free performance any more

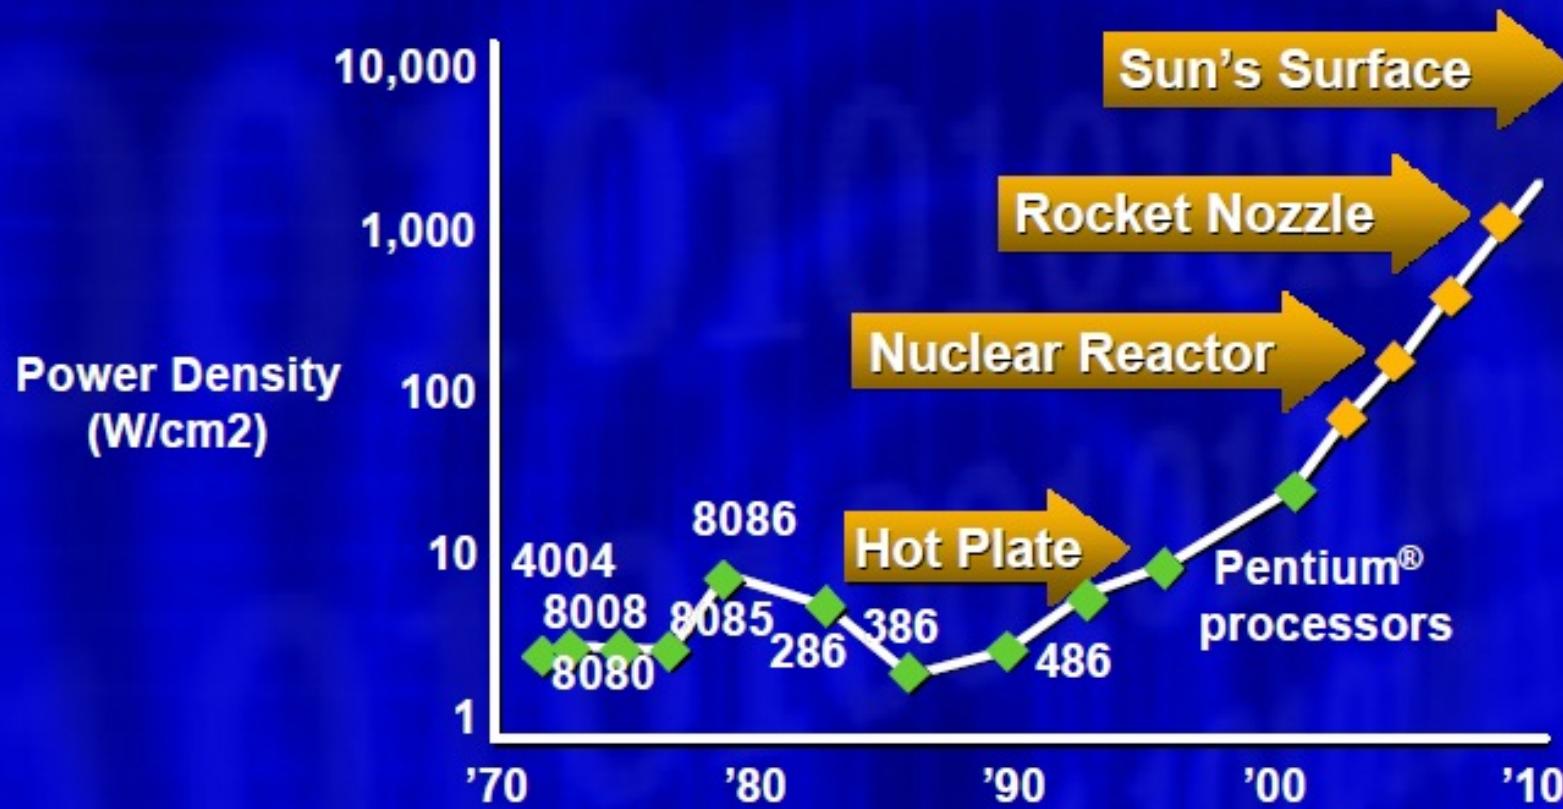


Performance is the **currency** of computing. You can often "buy" needed properties with performance.

Technology Scaling



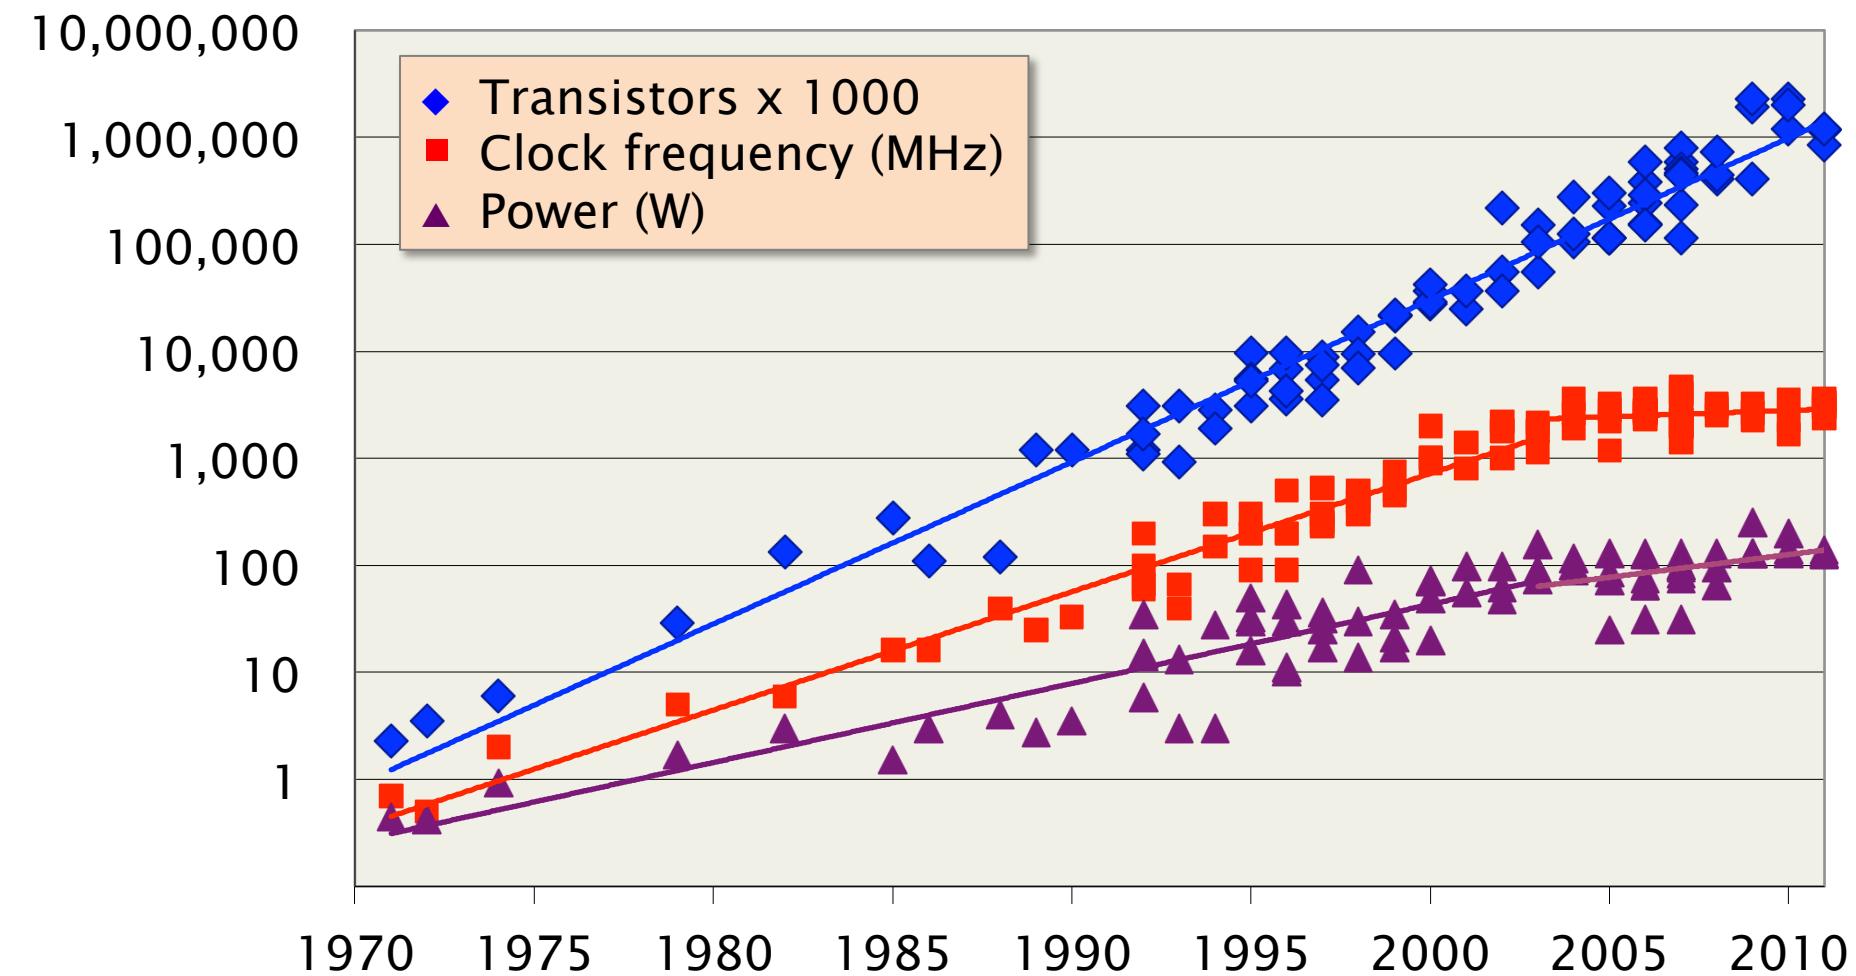
Power Density



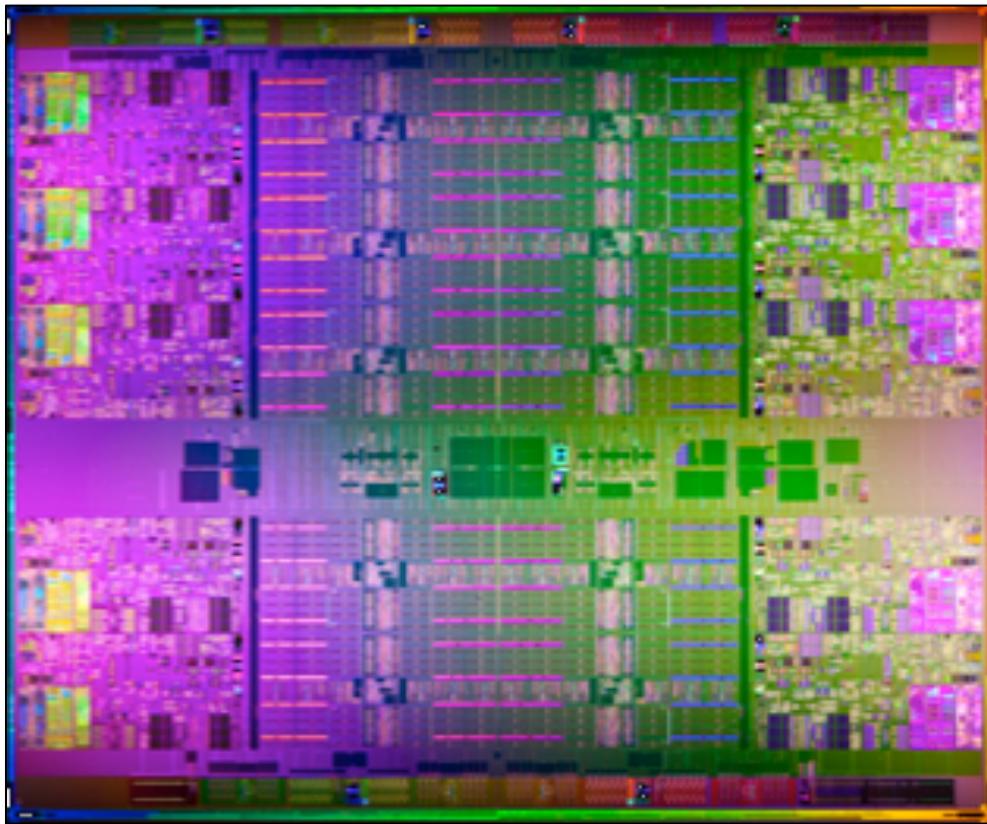
Source: Patrick Gelsinger, *Intel Developer's Forum*, Intel Corporation, 2004.

Power density, had scaling of clock frequency continued its trend of 25%–30% increase per year.

Technology Scaling



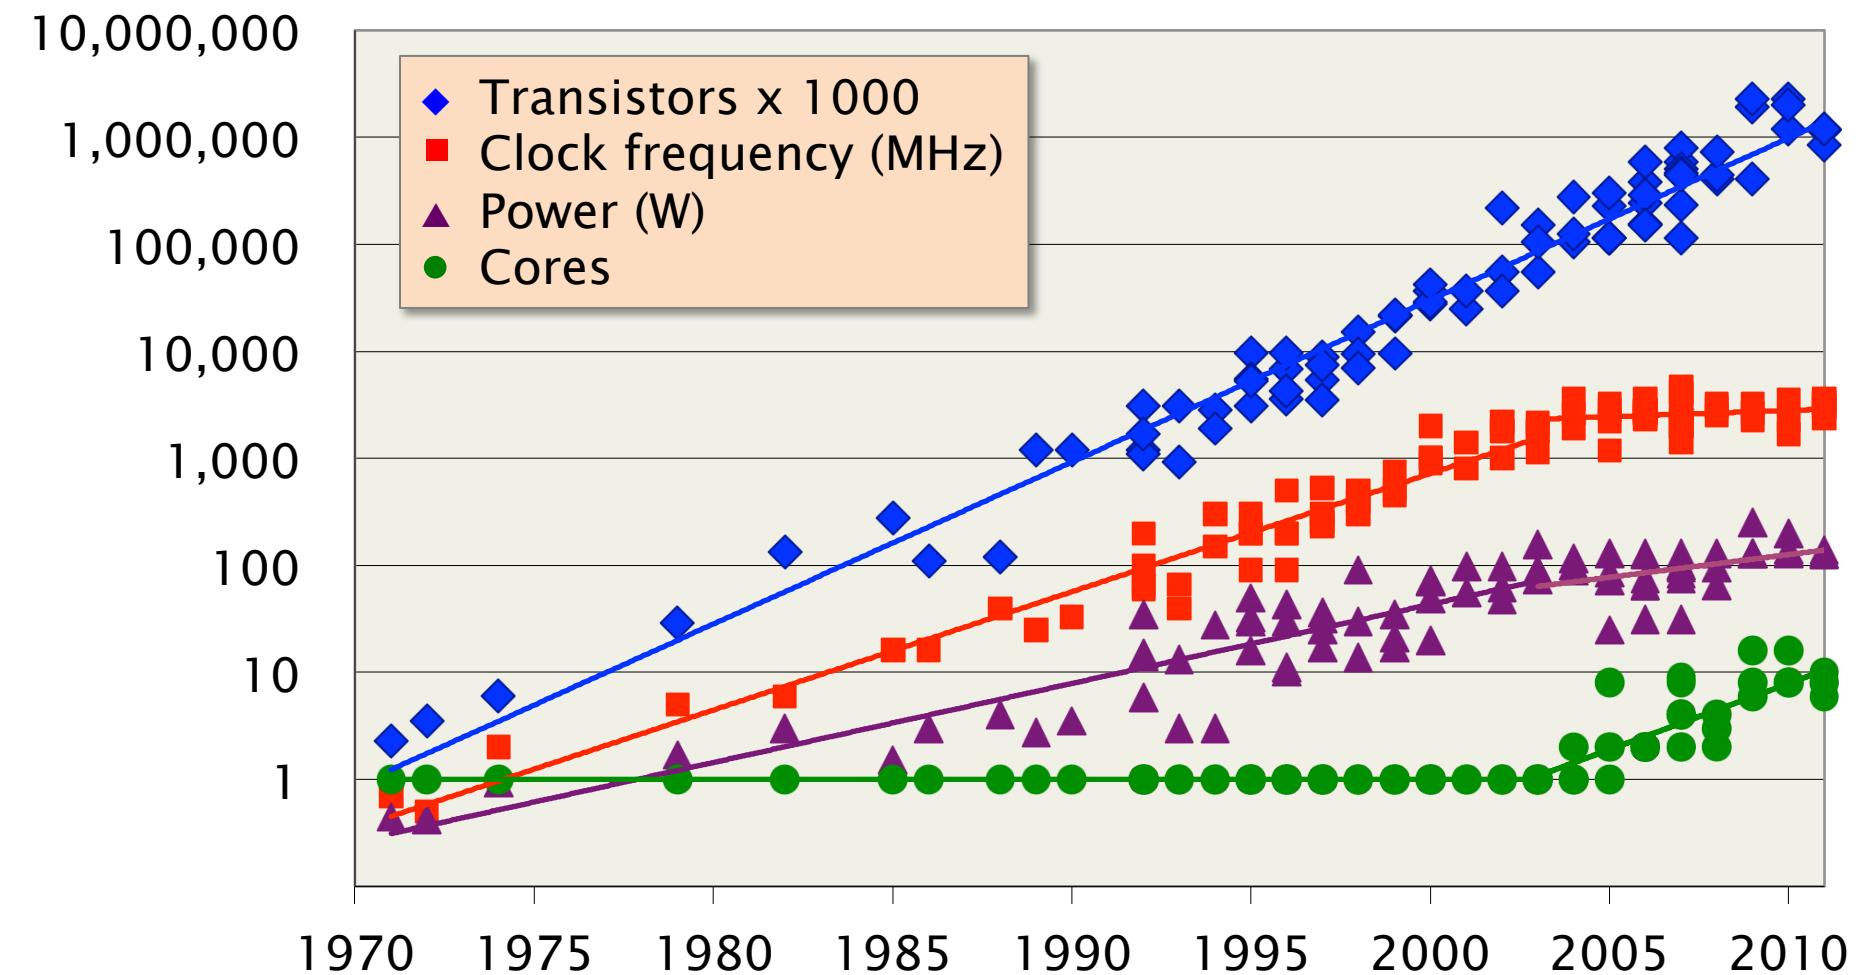
Vendor Solution: Multicore



Intel Xeon E7
(10 cores per chip)
Server systems
contain 4 chips

- To scale performance, put many processing cores on the microprocessor chip.
- Each generation of Moore's Law potentially doubles the number of cores.

Technology Scaling



Modern Era of Computing

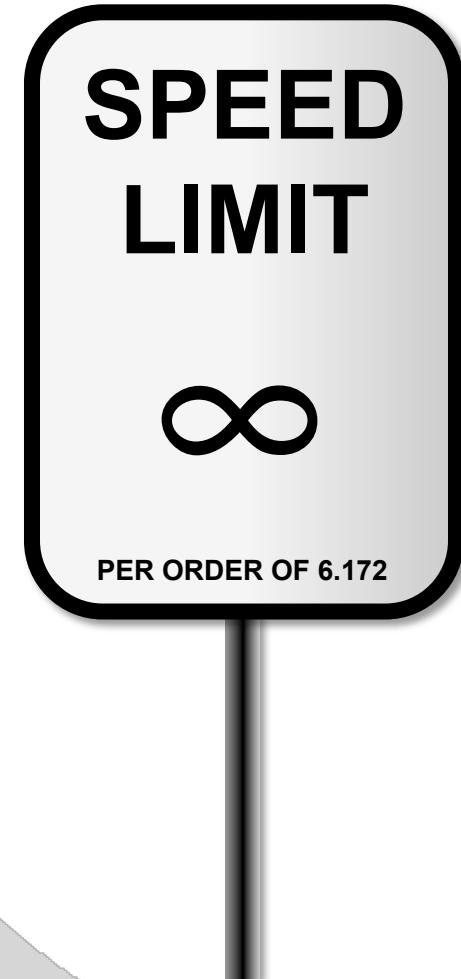
Moore's law is not giving free performance any more

Performance Engineering is the only way to get more currency



Performance is the **currency** of computing. You can often "buy" needed properties with performance.

6.172 ADMINISTRIVIA



Staff

Lecturers

- Prof. Charles E. Leiserson
- Prof. Saman P. Amarasinghe
- Dr. Aparna Chandramowlishwaran

Teaching assistants

- William Mitchell Leiserson
- Tao Benjamin Schardl
- Cong Yan
- Michael J Xu
- Damon Doucet

Administrative support

- Mary McDavitt
- Cree Bruins

Masters in the Practice of Software Systems Engineering (MITPOSSE)

- Expert programmers from industry who will review your code and provide feedback

Communication

Class home page

- <http://stellar.mit.edu/S/course/6/fa14/6.172/>

Correspondence

- <http://www.piazza.com/>
- All course material, project related, and administrative questions
- Mark personal communications to staff as private, but try to make most communications public

Recitations

“Learning the life skills needed
to be a true hacker”

Organization

- Two-hour duration
- Once a week on Friday
- “Lecture–studio” format
- Mandatory
- Must complete a set of tasks and get it checked off by your TA

Required Work

12% — Weekly Homeworks

30% — 2 Quizzes

- In class
- Closed book, but crib sheet allowed
- Details forthcoming

36% — Projects 1–3

- Beta, MITPOSSE review, final

22% — project 4

- Beta-1, MITPOSSE review, Beta2, Final

No final exam

Homework 1 & Recitation 1

Homework 1 is out today

- Part 1 is due by today @ 8:00PM!!
- Rest due on Monday

There will be a recitation tomorrow

- Go to any one of 3 recitations

TA Office Hours

Times

- Sunday, Monday, Tuesday, Wednesday
- 4:00 P.M. to 7:00 P.M.
- Location announced in lecture

Bring your laptop

- debugging support
- help with tools
- answers to conceptual questions
- good place to work on the projects with your team

Expectations

Required work is required!

- If you fail to do all the assignments, you risk failing the class
- If you miss a recitation assignment, you risk failing the class
- No late homework — hand in what you have by the deadline for partial credit

If an issue arises, please talk to your TA as soon as possible so that arrangements can be made!

Projects

You are given a correct but inefficient program.
Your mission: Make the program run fast.

Do whatever you can within the rules

- There is no right answer!
- Take advantage of the machine resources.
- Lots of creative freedom to explore many possible directions.
- Hard to be fastest, but easy to test which is fastest!

The journey is as important as the outcome

- You may try many things that will not give a performance improvement.
- Failure is as important as success → feedback!
- Tell us everything you did and why in your write-up.

Project Process

Project starts

Beta submission

- The staff will publish the performance results and a baseline for the final submission.
- Code from other projects will be available to all
 - Study and understand what they did, get inspired....but don't copy!

MITPOSSE design review

- After the Beta, you have just over a week to meet in a 90-minute design-review meeting with your assigned Master.
- Your Master will provide feedback on your code and design.
- Your Master will not grade you, but your attendance at the design review is mandatory.

Final submission

- Update the code to reflect your Master's comments.
- Enhance the performance to reach the published baseline.
 - Better than baseline → full credit for performance
 - Worse than baseline → fraction relative to the slowdown

Practicing engineers from the industry

- Unpaid volunteers who are contributing their time to help you!
- Senior engineers with lots of experience.
- You can learn a lot from them!

Please accord them proper respect

- Be responsive when they contact you to schedule the design review.
- Thank them for their feedback.
- Be personable.

What to expect

- Input on your code style
- Advice on deployment, testing debugging process
- What programmers do in industry
- Even career advice and mentoring

NOT how to make your program faster

Academic Honesty

Homework

- No sharing

While a project is active, you may share

- with your group

After beta, you will have access to everyone's code

- But... do not copy, just learn from them
- Practically, after looking at someone else's code, wait sometime before you touch your code (ex: one hour)

Use of outside materials

- You may use outside materials as long as you properly cite them.

Read the course information handout

- If you have any questions, please talk to your TA.

We will be using technology to detect cheating

Programming Languages

We will be using C

- Close to the metal
- Machine's memory is directly exposed
 - malloc() and free(), pointers, native data types
- Code compiles directly to machine language
- No hidden work (garbage collection, bounds checking)

Resources available on the class home page

- Manuals for various tools
- Quick references
- C primer on Monday the 8th @ 7:00PM

Online resources

- www.cprogramming.com
- search will find many other resources

Multicore Machine Resources

Cloud machines

- Collection of 12-core machines donated by Dell and Intel (thanks!)
- Log onto `cloud#.csail.mit.edu` for $\# = 0, 1, \dots, 4$
 - Need a CSAIL account to login
- Use by everyone (you need to load balance)
- For editing, compiling the code and running the tool frontends

Lanka machines

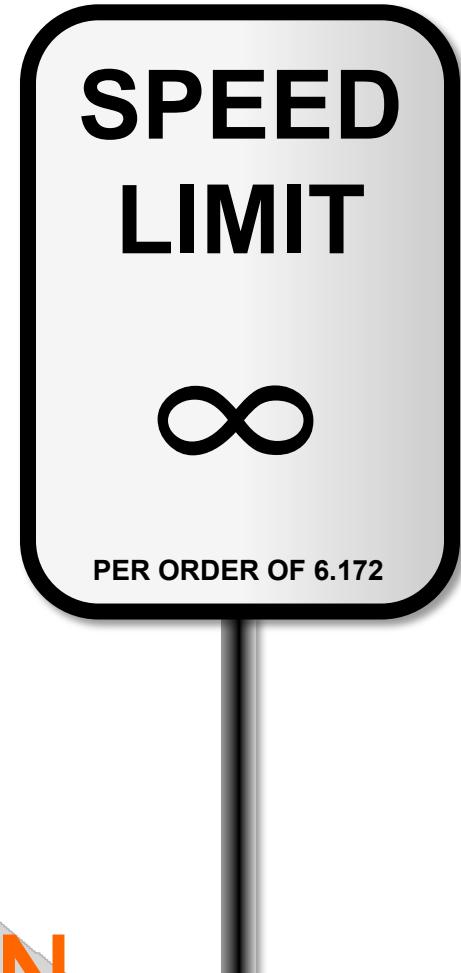
- Place to run the executables
- Only one job will be run on a machine at a time via a queue
- For performance number gathering

Intel E5-2695 V2 @ 2.4Ghz
12-cores (2 per node)
128 GB of DDR3 memory

Laptop development

- Recommended, but use at your own risk
- 6.172 staff will not help maintain your software

CASE STUDY: MATRIX MULTIPLICATION



Square-Matrix Multiplication

$$\begin{pmatrix} c_{11} & c_{12} & \cdots & c_{1n} \\ c_{21} & c_{22} & \cdots & c_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ c_{n1} & c_{n2} & \cdots & c_{nn} \end{pmatrix} = \begin{pmatrix} a_{11} & a_{12} & \cdots & a_{1n} \\ a_{21} & a_{22} & \cdots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{n1} & a_{n2} & \cdots & a_{nn} \end{pmatrix} \cdot \begin{pmatrix} b_{11} & b_{12} & \cdots & b_{1n} \\ b_{21} & b_{22} & \cdots & b_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ b_{n1} & b_{n2} & \cdots & b_{nn} \end{pmatrix}$$

C A B

$$c_{ij} = \sum_{k=1}^n a_{ik} b_{kj}$$

Assume for simplicity that $n = 2^k$.

Intel Xeon Computer System

Feature	Specification
Microarchitecture	Sandy Bridge
Clock frequency	2.4 GHz
Processor chips	2
Processing cores	8 per processor chip
Hyperthreading	2 way
Floating-point unit	8 double-precision operations per core per cycle
Cache-line size	64 B
L1-icache	32 KB private 8-way set associative
L1-dcache	32 KB private 8-way set associative
L2-cache	256 KB private 8-way set associative
L3-cache	20 MB shared 20-way set associative
DRAM	32 GB

$$\text{Peak} = 2 \times 8 \times 8 \times 2.4 \times 10^9 = 307 \text{ GFLOPS}$$

1. Triply Nested Loops in Python

```
import sys, random
from time import *

n = 4096

A = [[1.0*random.random()
      for row in xrange(n)]
      for col in xrange(n)]

B =
C =
start
for
end

print '%0.6f' % (end - start)
```

Running time
= 34,962 seconds
 \approx 9.75 hours

Is this fast?

Back-of-the-envelope calculation

$2n^3 = 2^{37}$ floating-point operations

Running time $\approx 2^{15}$ seconds

\therefore Python gets $2^{37}/2^{15} = 2^{22} \approx 4$ MFLOPS

Peak = 307 GFLOPS

Python gets $\approx 0.0013\%$ of peak

2. Let's Try Java

```
import java.util.Random;

public class mm_java {
    static int n = 4096;
    static double[][] A = new double[n][n];
    static double[][] B = new double[n][n];
    static double[][] C = new double[n][n];

    public static void main(String[] args) {
        Random r = new Random();

        for (int i=0; i<n; i++) {
            for (int j=0; j<n; j++) {
                A[i][j] = r.nextDouble();
                B[i][j] = r.nextDouble();
                C[i][j] = 0;
            }
        }

        long start = System.nanoTime();

        for (int i=0; i<n; i++) {
            for (int j=0; j<n; j++) {
                for (int k=0; k<n; k++) {
                    C[i][j] += A[i][k] * B[k][j];
                }
            }
        }

        long stop = System.nanoTime();

        double tdiff = (stop - start) * 1e-9;
        System.out.println(tdiff);
    }
}
```

Running time = 2,531 seconds
≈ 42 minutes
... about 14× faster than Python!
Still only 0.0177% of peak.

```
for (int i=0; i<n; ++i) {
    for (int j=0; j<n; ++j) {
        for (int k=0; k<n; ++k) {
            C[i][j] += A[i][k] * B[k][j];
        }
    }
}
```

3. Why Not C?

```
#include <stdlib.h>
#include <stdio.h>
#include <sys/time.h>
#include <assert.h>

typedef unsigned long long uint64_t;

#define n 4096
double A[n][n];
double B[n][n];
double C[n][n];

float tdiff (struct timeval *start,
             struct timeval *end) {
    return (end->tv_sec-start->tv_sec)
        + 1e-6*(end->tv_usec-start->tv_usec);
}

int main(int argc, const char *argv[]) {
    for (int i=0; i<n; ++i) {
        for (int j=0; j<n; ++j) {
            A[i][j] = (double)rand() / (double)RAND_MAX;
            B[i][j] = (double)rand() / (double)RAND_MAX;
            C[i][j] = 0;
        }
    }

    struct timeval start, end;
    gettimeofday(&start, NULL);

    for (int i=0; i<n; ++i) {
        for (int j=0; j<n; ++j) {
            for (int k=0; k<n; ++k) {
                C[i][j] += A[i][k] * B[k][j];
            }
        }
    }

    gettimeofday(&end, NULL);
    printf("%0.6f\n", tdiff(&start, &end));
    return 0;
}
```

Using the GCC compiler
Running time = 1,463 seconds
 \approx 24 minutes
... about 1.7 \times faster than Java.

```
for (int i=0; i<n; ++i) {
    for (int j=0; j<n; ++j) {
        for (int k=0; k<n; ++k) {
            C[i][j] += A[i][k] * B[k][j];
        }
    }
}
```

Where We Stand So Far

Version	Implementation	Time (s)	GFLOPS	Absolute speedup	Relative speedup	Fraction of peak
1	Python	34,962.21	0.004	1		0.00%
2	Java	2,530.65	0.054	14	13.8	0.00%
3	C, using GCC	1,462.50	0.094	24	1.7	0.00%

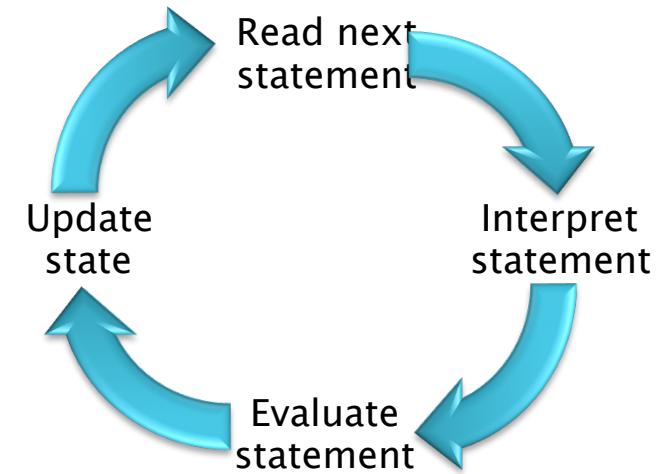
Why is Python so slow and C so fast?

- Python is interpreted.
- Java is compiled to byte-code, which is then interpreted and just-in-time (JIT) compiled.
- C is compiled directly to machine code.

Interpreter overhead

Lot of extra work

- Repeated again and again



JIT Compiler

- Do a lot of work once, and reuse
- Granularity of basic blocks or methods
 - basic block = instruction sequences without control flow
- When a new instruction/method, check if it is already jitted
- If already jitted, directly execute it
- If not, run the read/interpret/evaluate/update cycle and create a executable list of machine instructions

4. Optimization Switches

GCC provides a collection of optimization switches. Without touching the C code, we can just specify a switch to the compiler to ask it to optimize.

Opt. level	Meaning	Time (s)
-00	Do not optimize	1463
-01	Optimize	856
-02	Optimize even more	851
-03	Optimize yet more	427

4. Optimization Switches

Version	Implementation	Time (s)	GFLOPS	Absolute speedup	Relative speedup	Fraction of peak
1	Python	34,962.21	0.004	1	1	0.00%
2	Java	2,530.65	0.054	14	13.8	0.00%
3	C, using GCC	1,462.50	0.094	24	1.7	0.00%
4	+ switches	426.79	0.322	82	3.4	0.10%

How is our processor doing?

- We know only 0.1% of the peak
- But, why?
- What may be contributing to this slowdown?

Performance Counters

Modern hardware counts “events”

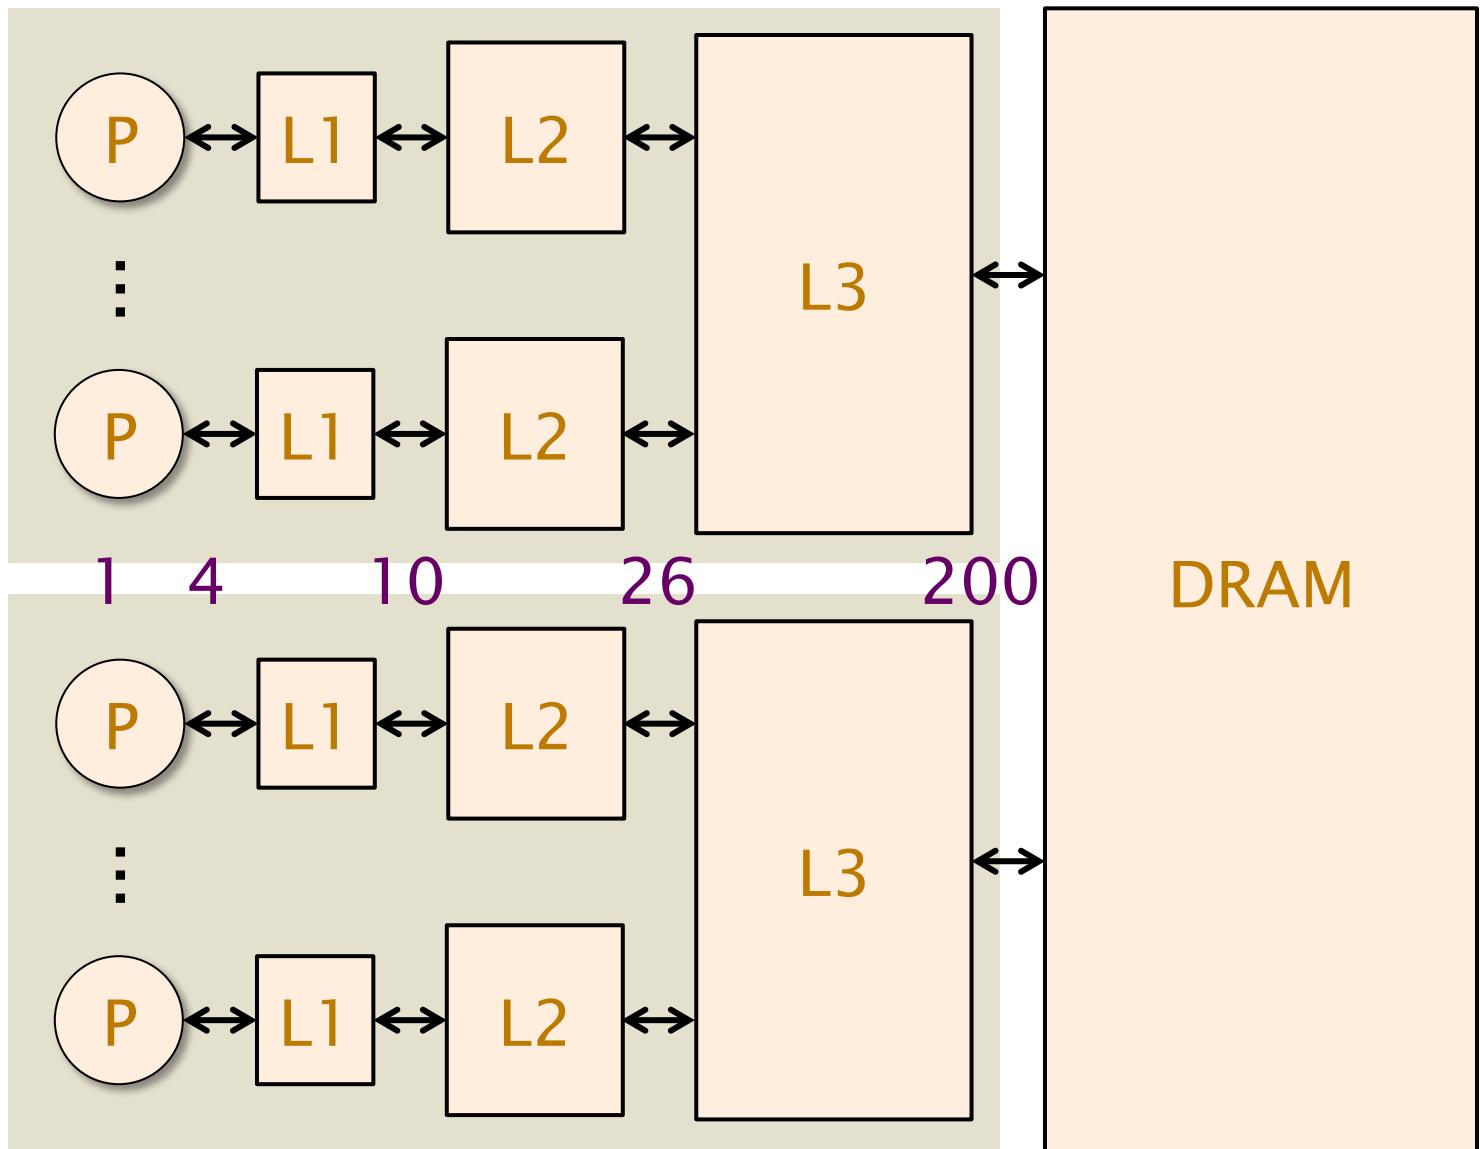
- Lot more information than just execution time

Version	Implementation	Run Time(s)	Cache References	Cache misses	Hit rate
4	+ switches	426.79	34,320,418,733	34,042,409,392	0.81%

Reference and misses for L3 cache

“Sandy Bridge” Memory Hierarchy

Latency
in clock
cycles



Performance Counters

Modern hardware counts “events”

- Lot more information than just execution time

Version	Implementation	Run Time(s)	Cache References	Cache misses	Hit rate
4	+ switches	426.79	34,320,418,733	34,042,409,392	0.81%

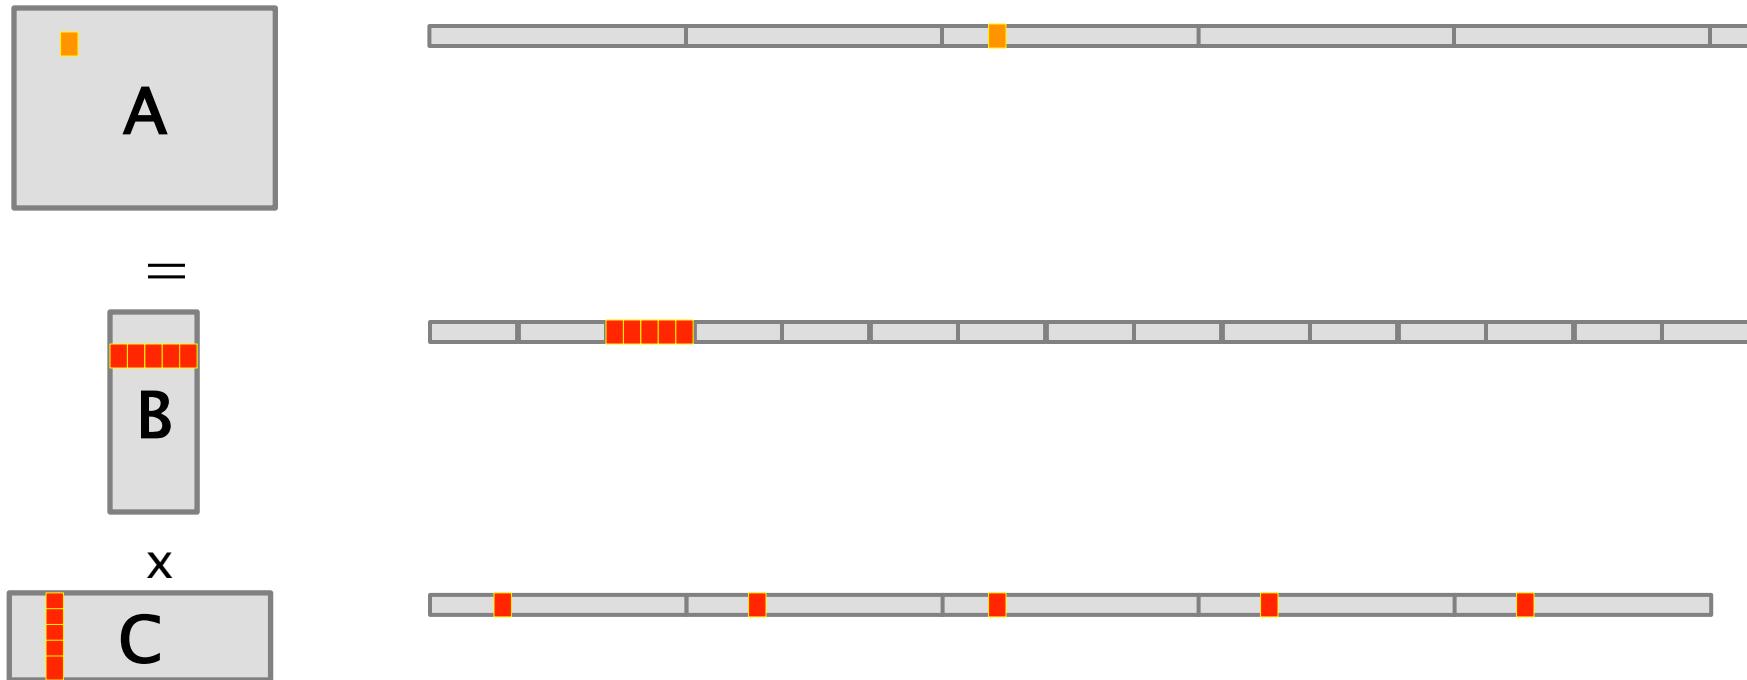
Reference and misses for L3 cache

Really bad hit rate

- should be in the 90's

5. Data Transpose

Scanning the memory



Contiguous accesses are better

- Data fetch as cache line (Core 2 Duo 64 byte L2 Cache line)
- Contiguous data → Single cache fetch supports 8 reads of doubles

Preprocessing of Data

In Matrix Multiply

- n^3 computation
- n^2 data

Possibility of preprocessing data before computation

- n^2 data $\rightarrow n^2$ processing
- Can make the n^3 happens faster

One matrix don't have good cache behavior

Transpose that matrix

- n^2 operations
- Will make the main matrix multiply loop run faster

5. Data Transpose

After transposing the C matrix

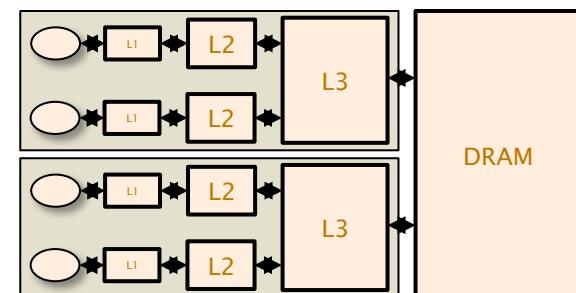
Version	Implementation	Time (s)	GFLOPS	Absolute speedup	Relative speedup	Fraction of peak
1	Python	34,962.21	0.004	1	1	0.00%
2	Java	2,530.65	0.054	14	13.8	0.00%
3	C, using GCC	1,462.50	0.094	24	1.7	0.00%
4	+ switches	426.79	0.322	82	3.4	0.10%
5	+ transpose	79.84	1.722	438	5.4	0.56%

Version	Implementation	Run Time(s)	Cache References	Cache misses	Hit rate
4	+ switches	426.79	34,320,418,733	34,042,409,392	0.81%
5	+ transpose	79.84	4,699,663,602	1,863,550,316	60.35%

~8x reduction of cache references!

- Why?

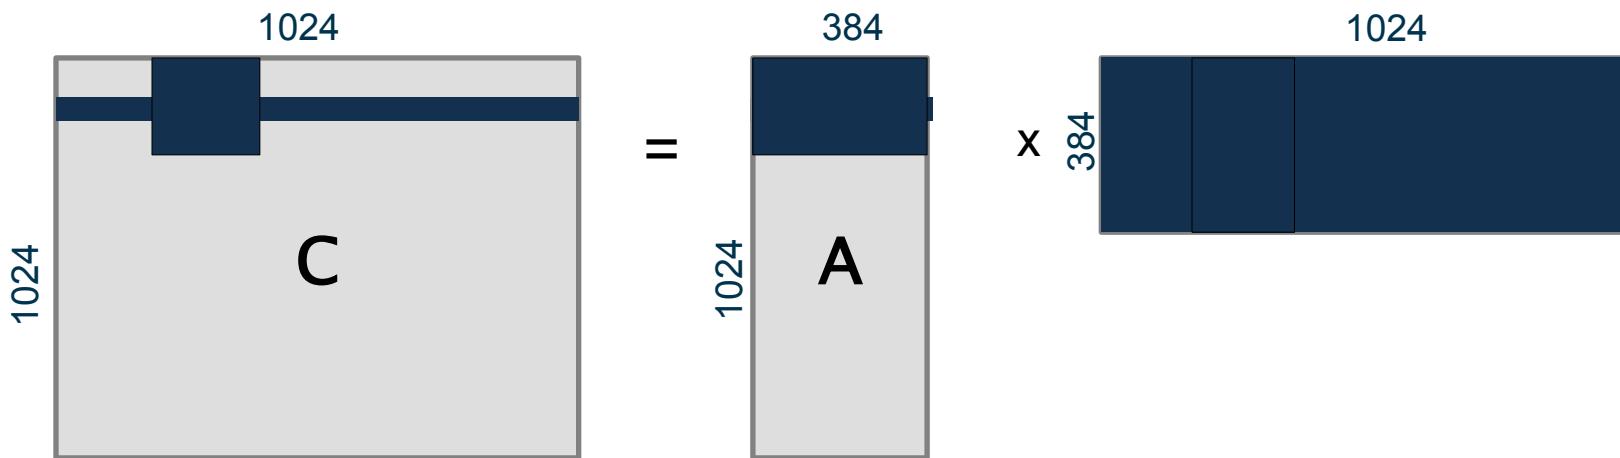
Improved hit rate



Data Reuse

Data reuse

- Change of computation order can reduce the # of loads to cache
- Calculating a row (1024 values of C)
 - C: $1024 \times 1 = 1024$ + A: $384 \times 1 = 384$ + B: $1024 \times 384 = 393,216$
= 394,524
- Blocked Matrix Multiply ($32^2 = 1024$ values of C)
 - C: $32 \times 32 = 1024$ + A: $384 \times 32 = 12,288$ + B: $32 \times 384 = 12,288$
= 25,600



Data Reuse

Data reuse

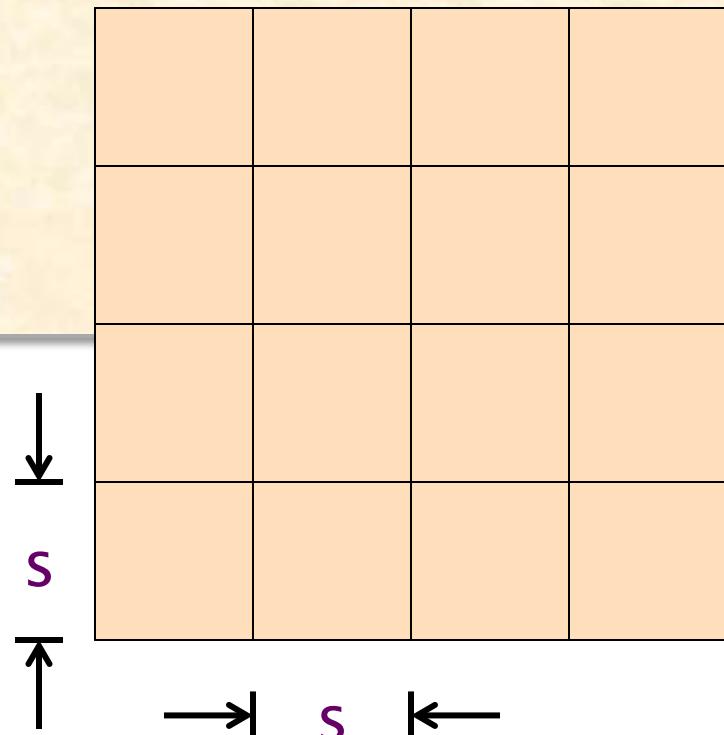
- Change of computation order can reduce the # of loads to cache
- Calculating a row (1024 values of C)
 - C: $1024 \times 1 = 1024$ + A: $384 \times 1 = 394$ + B: $1024 \times 384 = 393,216$
= 394,524
- Blocked Matrix Multiply ($32^2 = 1024$ values of C)
 - C: $32 \times 32 = 1024$ + A: $384 \times 32 = 12,284$ + B: $32 \times 384 = 12,284$
= 25,600
- If the data needed does not fit in the cache, has to be fetched each time needed from the next level
 - In calculating a Raw, each element used from B is reused only after 394,525 accesses to data (for the next raw)
 - In Blocked matrix multiply, each element used from B is reused after 65 accesses (and reused 32 times)
 - ◆ If the cache has more than 65 elements → cache hit!

Tiling

```
for (int ih = 0; ih < n; ih += s) {  
    for (int jh = 0; jh < n; jh += s) {  
        for (int kh = 0; kh < n; kh += s) {  
            for (int il = 0; il < s; il++) {  
                for (int jl = 0; jl < s; jl++) {  
                    for (int kl = 0; kl < s; ++kl) {  
                        C[ih+il][jh+jl] += A[ih+il][kh+kl] * B[kh+kl][jh+jl];  
                    }  
                }  
            }  
        }  
    }  
}
```

Cache misses

- If s^2 is sufficiently smaller than the size of the cache, the tiled loops incur only $\Theta(n^3/s)$ cache misses.



Performance of Tiling

Tile size	1 core (s)	16 cores (s)
1	859.64	65.30
2	309.60	26.32
4	178.17	9.19
8	93.14	6.56
16	76.65	5.09
32	60.85	3.49
64	54.17	3.24
128	44.79	2.76
256	40.87	4.73
512	42.77	6.79
1024	69.00	8.39
2048	65.96	26.13

Tile size

- For 16 cores, a 128×128 tile gives the best performance.
- For 1 core, however, 256×256 tile works best.
- If tile size is not properly tuned
 - either too large or too small
 - the code may perform poorly.
- Tiling is **fragile**.
 - Optimal size depends on many factors such as memory size and processor type
 - Need to retune to keep the edge

6. Divide and Conquer

IDEA: Tile for **every** power of 2.

$$\begin{bmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{bmatrix} = \begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix} \cdot \begin{bmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{bmatrix}$$
$$= \begin{bmatrix} A_{11}B_{11} & A_{11}B_{12} \\ A_{21}B_{11} & A_{21}B_{12} \end{bmatrix} + \begin{bmatrix} A_{12}B_{21} & A_{12}B_{22} \\ A_{22}B_{21} & A_{22}B_{22} \end{bmatrix}$$

- 8 multiplications of $n/2 \times n/2$ matrices.
- 1 addition of $n \times n$ matrices.

6. Divide and Conquer

```
void mmdac (double *C, double *A, double *B, int size) {  
    if (size <= 1) {  
        *C += *A * *B;  
    } else {  
        int s11 = 0;  
        int s12 = size/2;  
        int s21 = (size/2)*n;  
        int s22 = (size/2)*(n+1);  
        mmdac(C+s11, A+s11, B+s11, size/2);  
        mmdac(C+s12, A+s11, B+s12, size/2);  
        mmdac(C+s21, A+s21, B+s11, size/2);  
        mmdac(C+s22, A+s21, B+s12, size/2);  
        mmdac(C+s11, A+s12, B+s21, size/2);  
        mmdac(C+s12, A+s12, B+s22, size/2);  
        mmdac(C+s21, A+s22, B+s21, size/2);  
        mmdac(C+s22, A+s22, B+s22, size/2);  
    }  
}
```

Performance of D&C

Version	Implementation	Time (s)	GFLOPS	Absolute speedup	Relative speedup	Fraction of peak
1	Python	34,962.21	0.004	1	1	0.00%
2	Java	2,530.65	0.054	14	13.8	0.00%
3	C, using GCC	1,462.50	0.094	24	1.7	0.00%
4	+ switches	426.79	0.322	82	3.4	0.10%
5	+ transpose	79.84	1.722	438	5.4	0.56%
6	Dvidie-and-conquer	314.73	0.437	111	0.3	0.14%

Uh, oh! A big step backwards!

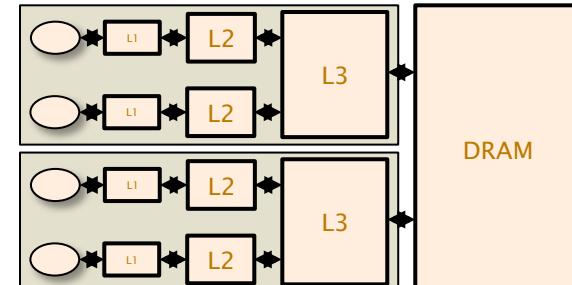
Version	Implementation	Run Time(s)	Cache References	Cache misses	Hit rate
4	+ switches	426.79	34,320,418,733	34,042,409,392	0.81%
5	+ transpose	79.84	4,699,663,602	1,863,550,316	60.35%
6	Dvidie-and-conquer	314.73	7,909,253,112	23,814,681	99.70%

Cache references nearly doubled

- Why?

But... we improved the hit rate

- Why?



Function-Call Overhead

```
void mmdac (double *C, double *A, double *B, int size) {  
    if (size <= 1) {  
        *C += *A * *B;  
    } else {  
        int s11 = 0;  
        int s12 = size/2;  
        int s21 = (size/2)*n;  
        int s22 = (size/2)*(n+1);  
        mmdac(C+s11, A+s11, B+s11, size/2);  
        mmdac(C+s12, A+s11, B+s12, size/2);  
        mmdac(C+s21, A+s21, B+s11, size/2);  
        mmdac(C+s22, A+s21, B+s12, size/2);  
  
        mmdac(C+s11, A+s12, B+s21, size/2);  
        mmdac(C+s12, A+s12, B+s22, size/2);  
        mmdac(C+s21, A+s22, B+s21, size/2);  
        mmdac(C+s22, A+s22, B+s22, size/2);  
    }  
}
```

The base case is too small.
We must **coarsen** the
recursion to avoid
function-call overhead.

7. Coarsening

```
void mmdac (double *C, double *A, double *B, int size) {  
    if (size <= CUTOFF) {  
        *C += *A * *B;  
    } else {  
        int s11 = 0;  
        int s12 = size/2;  
        int s21 = (size/2)*n;  
        int s22 = (size/2)*(n+1);  
        mmdac(C+s11, A+s11, B+s11, size);  
        mmdac(C+s12, A+s11, B+s12, size);  
        mmdac(C+s21, A+s21, B+s11, size);  
        mmdac(C+s22, A+s21, B+s12, size);  
  
        mmdac(C+s11, A+s12, B+s21, size);  
        mmdac(C+s12, A+s12, B+s22, size);  
        mmdac(C+s21, A+s22, B+s21, size);  
        mmdac(C+s22, A+s22, B+s22, size);  
    }  
}
```

The base case is too small. We must **coarsen** the recursion to avoid function-call overhead.

What is the CUTOFF value?

It depends...

Later we will learn about autotuning.

Performance of Coarsening + Transpose

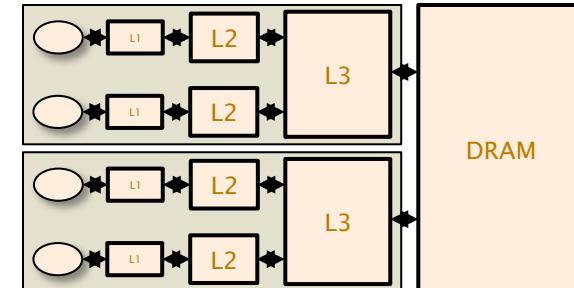
Version	Implementation	Time (s)	GFLOPS	Absolute speedup	Relative speedup	Fraction of peak
1	Python	34,962.21	0.004	1	1	0.00%
2	Java	2,530.65	0.054	14	13.8	0.00%
3	C, using GCC	1,462.50	0.094	24	1.7	0.00%
4	+ switches	426.79	0.322	82	3.4	0.10%
5	+ transpose	79.84	1.722	438	5.4	0.56%
6	Dvidie-and-conquer	314.73	0.437	111	0.3	0.14%
7	+ coarsening, transpose	79.04	1.739	442	4.0	0.57%

OK, we are back!

Version	Implementation	Run Time(s)	Cache References	Cache misses	Hit rate
4	+ switches	426.79	34,320,418,733	34,042,409,392	0.81%
5	+ transpose	79.84	4,699,663,602	1,863,550,316	60.35%
6	Dvidie-and-conquer	314.73	7,909,253,112	23,814,681	99.70%
7	+ coarsening, transpose	79.04	1,137,603,224	69,632,598	93.88%

Cache references are down by $\sim 7 \times$

- Why?



8. Vectorization

Each core of our computer has **8 vector units** which can initiate **8 floating-point operations** on each cycle using a single **vector instruction**, as long as the operations are independent. Most compilers can be induced to produce a **vectorization report**:

```
$ gcc -O3 -std=c99 mm_c.c -o mm_c_gcc_03 -ftree-vectorizer-verbose=2
...
Vectorizing loop at mm_c.c:14
mm_c.c:14: note: LOOP VECTORIZED.
...
mm_c.c:3: note: vectorized 1 loops in function.
mm_c.c:3: note: Completely unroll loop 15 times
```

```
for (int i=0; i<n; ++i) {
    for (int j=0; j<n; ++j) {
        for (int k=0; k<n; ++k) {
            C[i][j] += A[i][k] * B[k][j];
        }
    }
}
```

Interchange these
two loops.

8. Vectorization

Version	Implementation	Time (s)	GFLOPS	Absolute speedup	Relative speedup	Fraction of peak
1	Python	34,962.21	0.004	1		0.00%
2	Java	2,530.65	0.054	14	13.8	0.00%
3	C, using GCC	1,462.50	0.094	24	1.7	0.00%
4	+ switches	426.79	0.322	82	3.4	0.10%
5	+ transpose	79.84	1.722	438	5.4	0.56%
6	Dvidie-and-conquer	314.73	0.437	111	0.3	0.14%
7	+ coarsening, transpose	79.04	1.739	442	4.0	0.57%
8	+ vectorization	20.55	6.689	1,702	3.9	2.18%

Close to 4x speedup

9. Parallel Loops

We're running on only one of our 16 cores, leaving 15 idle. Let's use all of them!

```
cilk_for (int i=0; i<n; ++i) {
    cilk_for (int j=0; j<n; ++j) {
        for (int k=0; k<n; ++k) {
            C[i][j] += A[i][k] * B[k][j];
        }
    }
}
```

The `cilk_for` keyword, which is supported by latest version of GCC, indicates that all the iterations of the loop may execute in parallel.

8. Recursive Parallel Matrix Multiply

```
void mmdac (double *C, double *A, double *B, int size) {  
    if (size <= CUTOFF) {  
        *C += *A * *B;  
    } else {  
        int s11 = 0;  
        int s12 = size/2;  
        int s21 = (size/2)*n;  
        int s22 = (size/2)*(n+1);  
        cilk_spawn mmdac(C+s11, A+s11, B+s11, size/2);  
        cilk_spawn mmdac(C+s12, A+s11, B+s12, size/2);  
        cilk_spawn mmdac(C+s21, A+s21, B+s11, size/2);  
        mmdac(C+s22, A+s21, B+s12, size/2);  
        cilk_sync;  
        cilk_spawn mmdac(C+s11, A+s12, B+s21, size/2);  
        cilk_spawn mmdac(C+s12, A+s12, B+s22, size/2);  
        cilk_spawn mmdac(C+s21, A+s22, B+s21, size/2);  
        mmdac(C+s22, A+s22, B+s22, size/2);  
        cilk_sync;  
    }  
}
```

The named *child* function may execute in parallel with the *parent* caller.

Control may not pass this point until all spawned children have returned.

Parallel-Loops Performance

Version	Implementation	Time (s)	GFLOPS	Absolute speedup	Relative speedup	Fraction of peak
1	Python	34,962.21	0.004	1		0.00%
2	Java	2,530.65	0.054	14	13.8	0.00%
3	C, using GCC	1,462.50	0.094	24	1.7	0.00%
4	+ switches	426.79	0.322	82	3.4	0.10%
5	+ transpose	79.84	1.722	438	5.4	0.56%
6	Dvidie-and-conquer	314.73	0.437	111	0.3	0.14%
7	+ coarsening, transpose	79.04	1.739	442	4.0	0.57%
8	+ vectorization	20.55	6.689	1,702	3.9	2.18%
9	Parallel divide-andconquer	1.42	96.666	24,590	14.5	31.47%

Running time

- 14.5x performance improvement when using 16 cores!
- Now up to 1/3 of the peak performance!

10 & 11. Unportable Performance

- Use the `-march=corei7-avx` GCC compiler switch to generate modern AVX vector instructions, but the code won't run on older machines.
- Use the `-ffast-math -mavx` GCC compiler switch, which generates multiple clones of the code, one of which uses the AVX instructions. A test is made at runtime as to which version to use.
 - Portable code, but the performance is not portable.
- Use compiler intrinsics (assembly-language directives) to access the AVX instructions directly. Highly non-portable, but great performance!

Final Reckoning

Version	Implementation	Time (s)	GFLOPS	Absolute speedup	Relative speedup	Fraction of peak
1	Python	34,962.21	0.004	1	1	0.00%
2	Java	2,530.65	0.054	14	13.8	0.00%
3	C, using GCC	1,462.50	0.094	2.1	1.7	0.00%
4	+ switches	426.79	0.322	82	3.4	0.10%
5	+ transpose	79.84	1.72	438	5.4	0.56%
6	Ddivide-and-conquer	314.73	0.437	111	0.3	0.14%
7	+ coarsening, transpose	79.04	0.739	442	4.0	0.57%
8	+ vectorization	20.55	6.689	1,702	3.9	2.18%
9	Parallel divide-andconquer	1.72	96.666	24,590	14.5	31.47%
10	+ machine-specific compilation	1.05	131.465	33,442	1.4	42.80%
11	+ AVX intrinsics	0.67	206.288	52,479	1.6	67.15%

52,479x Improvement!

Our sub-1-second Version 11 rivals the experts who coded the Intel Math Kernel Library!

Performance Engineering



Gas economy

45,978 ×



- You won't generally see the kind of performance improvement we obtained for matrix multiplication.
- But in 6.172 you will learn how to print the currency of performance all by yourself.